

Notes and Tutorial on GDB

Jack Rosenthal

CSM Linux Users Group

Note

This isn't a complete reference, but an unprofessional handy source. You might find out a lot more by reading the manual.

Contents

1	Introduction	1
1.1	Compiling your programs for GDB	2
1.2	A basic example	3
1.3	The Terminal User Interface	5
2	Basic Features	5
2.1	Stepping through your code	5
2.2	Working with Variables	6
2.3	Breakpoints	8
2.4	Watchpoints	8
2.5	Disabling and Enabling {Break,Watch}points	8
3	Advanced Magical Wizardry	9
3.1	Calling Functions	9
3.2	Examining Memory	9
3.3	When things get so bad to the point where you are dealing with assembly code . . .	11

1 Introduction

The GNU Debugger (GDB) is an awesome tool for debugging your programs written in C, C++, FORTRAN (please don't), and plenty of other compiled languages. You can use it by running `gdb` from the command line.

1.1 Compiling your programs for GDB

GDB can **make use of special symbols** in your program to **help you debug**. Add the **-ggdb3** flag to **include these symbols**. Example below:

```
$ gcc -ggdb3 -std=c11 -o program.debug program.c
```

This compiler flag works with `gcc`, `clang`, or any of their related compilers (`g++`, etc.). Generally, you will only want to **use this flag for debugging**, as it not only embeds the source of your program in the binary but also increases the size of the binary significantly. The rest of this document assumes you have compiled with this flag when debugging.

Many people add a `debug` target to their `Makefile`, this allows them to compile their program for debugging and start the debugger simply typing `make debug`. The following listing contains a complete `Makefile` example, it may be overkill for small projects.

Listing 1: A complete Makefile example

```
CXX = clang  定义clang编译器
SRCFILES = $(wildcard src/*.c) 使用函数获取src目录下所有以.c结尾的文件列表
OBJFILES = $(patsubst src/%.c, bin/%.o, $(SRCFILES)) 将.c文件路径替换为对应的.o文件路径
OBJDEBUG = $(patsubst src/%.c, bin/%.debug.o, $(SRCFILES)) 同上
OUTFILE = myprogram 输出的二进制文件名称
CXXFLAGS = -Wall -std=c11 编译标志, 打开所有警告、使用C11标准
LDFLAGS = -lreadline 链接标志, 指定链接时需要使用readline库
          语法, 目标项: 目标所依赖项
          (TAB键) 生成目标相所执行的规则
all: $(OUTFILE)

$(OUTFILE): $(OBJFILES)
    $(CXX) $(LDFLAGS) $(OBJFILES) -o $(OUTFILE)

bin/%.o: src/%.c  将目标文件输出到文件名为$@的位置, $@表示当前目标的名称
    $(CXX) $(CXXFLAGS) -o $@ -c $<
    编译而不链接, 并指定编译的源文件为$<, $<表示依赖列表中的第一个文件

# 'make debug' depends on myprogram.debug. Make will check if this file
# exists and is up to date, and if so, call the debugger.
debug: $(OUTFILE).debug
    gdb -tui $(OUTFILE).debug
    调试时打开 text-user-interface

# 'make myprogram.debug' links the binary we are looking for.
$(OUTFILE).debug: $(OBJDEBUG)
    $(CXX) $(LDFLAGS) $(OBJDEBUG) -o $(OUTFILE).debug

# This makes each individual object with debugging symbols
bin/%.debug.o: src/%.c
    $(CXX) $(CXXFLAGS) -ggdb3 -o $@ -c $<
```

This `Makefile` **compiles a multifile project** with all of the source located at `src/*.c`.

1.2 A basic example

A coworker has handed you a simple C program they wrote, **but is segfaulting**. They are fairly new to C programming¹ and would like your help debugging. The program they wrote follows:

Listing 2: Your coworker's program (prog1.c)

```
1  #include <stdio.h>
2  #include <stdlib.h>
3
4  /**
5   * Allocate memory and read a string from standard input
6   */
7  char *
8  read_str_buf (size_t buf_sz)
9  {
10     char    *buf;
11     buf = malloc(buf_sz);
12     buf = fgets(buf, buf_sz - 1, stdin);
13     return buf;
14 }
15
16 int
17 main (argc, argv)
18     int    argc;
19     char    **argv;
20 {
21     char    *str;
22     for (;;) {
23         printf("What would you like me to print? ");
24         str = read_str_buf(1<<31);
25         printf("%s\n", str);
26     }
27     free(str);
28     return 0;
29 }
```

Can you spot the bug? Not easy without a debugger. Let's throw it in `gdb`! Start off by compiling with debugging symbols and starting `gdb` on the program.

```
$ gcc -ggdb3 -o prog1.debug prog1.c
$ gdb prog1.debug
```

¹Your coworker also used the 1st edition K&R C book to learn

When gdb first starts, you will see this:

```
GNU gdb (GDB) 7.10.1
Copyright (C) 2015 Free Software Foundation, Inc.
...
For help, type "help".
Type "apropos word" to search for commands related to "word"...
Reading symbols from prog1.debug...done.
(gdb)
```

gdb has near a gazillion commands, but we are going to want to use either `start` or `run` here. `start` will create a breakpoint at the start of `main` before running, whereas `run` will run right away. If we wanted the program to read a file on standard input, we could use the syntax `run < myfile` to read it. Go ahead and type `run`, and give it the input "Hello, World!":

```
(gdb) run
Starting program: prog1.debug
What would you like me to print? Hello, World!

Program received signal SIGSEGV, Segmentation fault.
0x00007ffff7aa0eb0 in __GI__IO_getline_info () from /usr/lib/libc.so.6
```

Yikes! What caused that? We can generate a backtrace to see where this happened. To do this, type `backtrace` (or, if you are lazy, `bt` works too). 这是一个堆栈，先进后出，所以最下面的行是先执行的代码

```
(gdb) bt
#0 0x00007ffff7aa0eb0 in __GI__IO_getline_info () from /usr/lib/libc.so.6
#1 0x00007ffff7a9fe1d in fgets () from /usr/lib/libc.so.6
#2 0x0000000000400650 in read_str_buf (buf_sz=...) at prog1.c:12
#3 0x0000000000400680 in main (argc=1, argv=0x7fffffff7e8) at prog1.c:24
```

Well, obviously it was at that `fgets` on line 12. Let's kill the program and restart it. We'll `break` at the start of `read_str_buf` to inspect what may be happening there. 执行了break后再执行run，程序会在break位置暂停

```
(gdb) break read_str_buf
Breakpoint 1 at 0x400622: file prog1.c, line 11.
(gdb) run
Starting program: prog1.debug

Breakpoint 1, read_str_buf (buf_sz=...) at prog1.c:11
11         buf = malloc(buf_sz);
```

Next, let's print some variables and step through some lines. We are going to use a few commands:

`print varname` Print the value of `varname`, can be abbreviated to `p`

`next` Go to the next line, can be abbreviated to `n`

```

(gdb) print buf_sz
$1 = 18446744071562067968
(gdb) next
12      fgets(buf, buf_sz, stdin);
(gdb) print buf
$2 = 0x0
(gdb) next
What would you like me to print? Hello, World?

Program received signal SIGSEGV, Segmentation fault.
0x00007ffff7aa0eb0 in __GI__IO_getline_info () from /usr/lib/libc.so.6

```

Hopefully you said “Hey that’s it, malloc returned NULL” when you printed buf. You can go ahead and type quit now, we will come back to this example later.

Brownie points are available if you can name why:

1. Even if the program worked, and allocated a reasonable amount of memory (say 1024), it would **still leak memory**.
2. The prompt line did not print until we reached the `fgets`
3. `buf_sz` was 18446744071562067968, not 2147483648
4. If line 12 was changed to read:

```
buf = fgets(buf, buf_sz, stdin);
```

the program would not segfault until line 25.

1.3 The Terminal User Interface

Yes. gdb has a nice **terminal user interface** that you can access by **starting gdb with the `-tui` flag**. This allows you to view each line in the source code you are stepping through in a nice box above the debugger. Sometimes the interface will get messed up and you will need to redraw, simply type **Ctrl+L** to redraw it.

2 Basic Features

2.1 **Stepping** through your code

For stepping through your code, you will want to use the **step, next, and continue** (abbreviated to `s`, `n`, and `c`) commands.

- The **step** command will execute whatever is on the current line **stepping into function calls** if there are any.

- The `next` command will execute whatever is on the current line *continuing through function calls* if there are any. `step`: 会跳进函数的首行, `next`: 会直接执行完函数, 转到下一行, `c`: 继续运行代码直到遇到断点
- The `continue` command will *keep running your code until it reaches a breakpoint*.

Using the program from earlier, here is an example. You may find it helpful to use the TUI, so you can see which line you are at visually.

```
(gdb) start
Temporary breakpoint 1 at 0x400669: file prog1.c, line 23.
Starting program: prog1.debug

Temporary breakpoint 1, main (argc=1, argv=0x7fffffff7e8) at prog1.c:23
23         printf("What would you like me to print? ");
(gdb) next
24         str = read_str_buf(1<<31);
(gdb) step
read_str_buf (buf_sz=18446744071562067968) at prog1.c:11
11         buf = malloc(buf_sz);
(gdb) next
12         buf = fgets(buf, buf_sz - 1, stdin);
(gdb) next
What would you like me to print? Hello, World?

Program received signal SIGSEGV, Segmentation fault.
0x00007ffff7aa0eb0 in __GI_IO_getline_info () from /usr/lib/libc.so.6
```

Notice how we used `step` to step into the `read_str_buf` function call. If we used `next`, it would silently run that function without taking us through it.

2.2 Working with Variables

`gdb` gives you full control of your variables at any time.

<code>print varname</code>	Print the value of <i>varname</i> , can be abbreviated to <code>p</code>
<code>set varname=value</code>	Set <i>varname</i> to <i>value</i>
<code>print varname=value</code>	Set <i>varname</i> to <i>value</i> and print it

value can be another variable.

To change frames, type `backtrace` to find the frame number to change to, then type `frame n` to change to frame *n*.

To list all variables in this frame, type `info locals`. To list the arguments to this frame, type `info args`. To get way too much information about this frame, type `info frame`.

Listing 3: prog2.c

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  extern char **environ; 将系统环境变量存储在一个全局指针数组变量environ中
4
5  /**
6   * My own strlen implementation
7   */
8  int my_strlen(const char *in) {
9      const char *p = in;
10     while (*p++); 先执行p++,再解引用,因为运算符++优先级高于解引用运算符*
11     return p - in - 1;
12 }
13
14 /**
15  * My own strdup implementation
16  * I think there's some sort of bug in this... but I can't find it!
17  */
18 char * my_strdup(const char *in) {
19     char *new = malloc(my_strlen(in) + 1), *p = new;
20     while (*in) {
21         *p = *in;
22         p++, in++;
23     } 这里有bug,*in在等于\0时,退出while循环,但是p的最后一个字符没有被初始化。*p = '\0';
24     return new;
25 }
26
27 int main(int argc, char **argv) {
28     char **p = environ;
29     while (*p) {
30         char *str = my_strdup(*p);
31         printf("%s\n", str);
32         free(str);
33         p++;
34     }
35     return 0;
36 }
```

2.3 Breakpoints

To set a **breakpoint**, use the incredibly powerful **break** command. The general syntax is:

```
break [(location)] [if (condition)]
```

where (location) can be:

A line number (current file assumed)	lineno
A file name and line number	file.c:lineno
A function name	my_function
A function name with namespace (for C++)	std::my_function
A function with types (for overloading in C++)	my_function(int)
The address of the program counter	*0xfedcba76

and (condition) can be any valid conditional syntax in your language (eg. `i == 32`). If the location is not specified, the current program counter is assumed.

即申请资源并拷贝字符串: `malloc + strcpy`

For our example, we are going to use the `prog2.c` program, which reimplements `strdup`, but with a bug. The main function prints all of the user's environment variables after they've been through `my_strdup` so you can see the bug (compare to the output of `env`).

Load up this program in `gdb` and issue the following commands. Keep pressing `c` whenever you reach a breakpoint, and notice how breakpoint 2 will only trigger when the first letter of the environment variable is H. Neato!

```
(gdb) break my_strlen
(gdb) break my_strdup if *in == 'H'
(gdb) run
```

2.4 Watchpoints

You may want to break the program whenever a variable is changed or read. This is called setting a **watchpoint**. The syntax is:

watch x break whenever x is changed.

rwatch x break whenever x is read.

命令使用示例放在PDF最后面了

awatch x break upon read or write

2.5 Disabling and Enabling {Break,Watch}points

If you are tired of breaking, you may want to disable a specific breakpoint. To list breakpoints and their associated number `n`, type **info breakpoints**. Then, type **disable n** or **enable n**.

```
(gdb) break my_strlen
Breakpoint 1 at 0x11b5: file prog2.c, line 6. // 我本地代码的行号
(gdb) break my_strdup
Breakpoint 2 at 0x11ee: file prog2.c, line 12.
(gdb) info breakpoints
Num    Type             Disp Enb Address            What
1      breakpoint       keep y   0x00000000000011b5 in my_strlen at prog2.c:6
2      breakpoint       keep y   0x00000000000011ee in my_strdup at prog2.c:12
(gdb) disable 2
(gdb) info b
Num    Type             Disp Enb Address            What
1      breakpoint       keep y   0x00000000000011b5 in my_strlen at prog2.c:6
2      breakpoint       keep n   0x00000000000011ee in my_strdup at prog2.c:12
(gdb)
```


3 Advanced Magical Wizardry

3.1 Calling Functions

`gdb` allows you to call one of your functions, or any function linked to your program. Use the `call` command. Here is an example:

```
(gdb) start
Temporary breakpoint 1 at 0x4006b7: file prog2.c, line 28.
Starting program: prog2.debug

Temporary breakpoint 1, main (argc=1, argv=0x7fffffffe7c8) at prog2.c:28
28      char **p = environ;
(gdb) call my_strlen("Hello, World!")
$1 = 13
```

3.2 Examining Memory

`gdb` has an amazing command, `x`, but few people actually remember how to use it without typing `help x`. Here is its syntax:

$$x/\langle\text{amount}\rangle\langle\text{format}\rangle\langle\text{size}\rangle\ \langle\text{variable or address}\rangle$$

`$\langle\text{amount}\rangle$` is the amount of objects of the specified size to print. It will default to 1 if you don't specify it.

`$\langle\text{format}\rangle$` is a single character that specifies how it should look when printed. Here are the different options:

- `o` octal
- `x` hexadecimal
- `d` decimal
- `u` unsigned decimal
- `t` binary
- `f` float
- `a` address
- `i` instruction
- `c` char
- `s` string
- `z` zero padded hexadecimal

`$\langle\text{size}\rangle$` specifies the size of the data to be printed. It is also a single character. Here are the different options:

- `b` byte (8 bits)
- `h` half word (16 bits)
- `w` word (32 bits)
- `g` giant word (64 bits)

That was awfully complex. For an example, let's examine `prog3.c`.

Listing 4: prog3.c

```
1  #define _POSIX_C_SOURCE 200809L
2  #include <stdio.h>
3  #include <string.h>
4
5  int main(int argc, char **argv) {
6      char *str = strdup("Hello, World!\n");
7      int x = 42;
8      int *y = &x;
9      double pi = 3.1415926535897932384;
10     double *pi_ptr = &pi;
11     if (*str && x && *y && *(int *)pi_ptr)
12         printf("Goodbye, world!\n");
13     return 0;
14 }
```

Here are some examples:

```
(gdb) start
Temporary breakpoint 1 at 0x400555: file prog3.c, line 6.
Starting program: prog3.debug

Temporary breakpoint 1, main (argc=1, argv=0x7fffffffefc8) at prog3.c:6
6      char *str = strdup("Hello, World!\n");
(gdb) n
7      int x = 42;
(gdb) x/s str
0x601010: "Hello, World!\n"
(gdb) x/c str
0x601010: 72 'H'
(gdb) x/6c str
0x601010: 72 'H' 101 'e' 108 'l' 108 'l' 111 'o' 44 ','
(gdb) n
8      int *y = &x;
(gdb) n
9      double pi = 3.1415926535897932384;
(gdb) x/tw y
0x7fffffffffffc4: 000000000000000000000000000000000101010
(gdb) n
10     double *pi_ptr = &pi;
(gdb) n
11     if (*str && x && *y && *(int *)pi_ptr)
(gdb) x/tg pi_ptr
0x7fffffffffffb8: 0100000000000100100100001111101101010100010001000010110100011000
```

3.3 When things get so bad to the point where you are dealing with assembly code

You may want to disassemble part of the program:

```
(gdb) disassemble main
Dump of assembler code for function main:
   0x000000000400546 <+0>:      push    %rbp
   0x000000000400547 <+1>:      mov     %rsp,%rbp
   0x00000000040054a <+4>:      sub     $0x40,%rsp
   0x00000000040054e <+8>:      mov     %edi,-0x34(%rbp)
   ...
```

Maybe examine the registers:

```
(gdb) info registers
rax            0x601010          6295568
rbx            0x0              0
rcx            0xa21646c726f57  2851464966991703
rdx            0x0              0
rsi            0x400657          4195927
...
```

You can also use `nexti` and `stepi` to step instructions (rather than lines).

You my friend, have achieved GDB wizard status.

```
(gdb) start
Temporary breakpoint 1 at 0x1237: file prog2.c, line 22.
Starting program: /home/xxx/prog2.debug
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib/x86_64-linux-gnu/libthread_db.so.1".
```

```
Temporary breakpoint 1, main (argc=1, argv=0x7fffffffe368) at prog2.c:22
```

```
22 char** p = environ;
```

```
(gdb) watch p // 观察p的改变
```

```
Hardware watchpoint 2: p
```

```
(gdb) next // p被赋值了, 即改变了
```

```
Hardware watchpoint 2: p
```

```
Old value = (char **) 0x0
```

```
New value = (char **) 0x7fffffffe378 // 改变前后的信息
```

```
main (argc=1, argv=0x7fffffffe368) at prog2.c:23
```

```
23 while (*p) {
```

```
(gdb) next
```

```
24 char* str = my_strdup(*p);
```

```
(gdb) next
```

```
25 printf("%s\n", str);
```

```
(gdb) next
```

```
SHELL=/bin/bash
```

```
26 free(str);
```

```
(gdb) next
```

```
27 p++; // p地址又改变了
```

```
(gdb) next
```

```
Hardware watchpoint 2: p
```

```
Old value = (char **) 0x7fffffffe378
```

```
New value = (char **) 0x7fffffffe380 // 改变前后的信息
```

```
main (argc=1, argv=0x7fffffffe368) at prog2.c:23
```

```
23 while (*p) {
```

```
(gdb)
```